

specific algorithms

W2. Discrete Classification I - Logistic Regression

Guang Cheng

University of California, Los Angeles

guangcheng@ucla.edu

Week 2

Classification

"averaged 0-1 loss" 是指对分类模型进行评估时使用的一种损失函数。在机器学习和统计学中，0-1损失函数是一种常用的损失函数，用于评估分类模型的性能。

0-1损失函数定义如下：如果预测正确，则损失为0，否则为1。

- A typical dataset in classification $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$.
 - \mathbf{x}_i : the covariate vector of i -th instance
 - $y_i \in \{-1, 1\}$: binary label of i -th instance
- **Question:** Can we directly minimize the averaged 0-1 loss?

$$\text{Training Error} : \frac{1}{n} \sum_{i=1}^n I(f(\mathbf{x}_i) \neq y_i)$$

指下凸函数，因为下凸才有最小值。

- **Answer:** No, the 0-1 loss function is non-convex and discontinuous, so (sub)gradient methods cannot be applied.

Classification - surrogate loss

替代损失函数试图近似0-1损失函数的行为，同时还要保持可导和凸性质，以便利用梯度下降法等优化算法。常见的替代损失函数包括Logistic损失、交叉熵损失等。

为计算非连续0-1损失下 loss function 的导数

- We can replace the 0-1 loss by other loss functions, say surrogate loss

$$\frac{1}{n} \sum_{i=1}^n I(f(\mathbf{x}_i) \neq y_i) \Rightarrow \frac{1}{n} \sum_{i=1}^n L(f(\mathbf{x}_i), y_i) = \frac{1}{n} \sum_{i=1}^n \phi(f(\mathbf{x}_i)y_i)$$

两种 surrogate loss 的类型:

- Hinge loss: $\phi(x) = \max\{0, 1 - x\}$
- Logistic loss $\phi(x) = \log(1 + \exp(-x))$

为计算非连续0-1损失


在整个实数域内可导

One surrogate loss - Hinge Loss

- Definition of Hinge loss:

↙ 衡量预测值与实际值之间的差异
 \hat{y} 与 y 之间的差异

$$L_{\text{hinge}}(f(\mathbf{x}_i), y_i) = \begin{cases} 1 - f(\mathbf{x}_i)y_i & \text{if } f(\mathbf{x}_i)y_i \leq 1 \\ 0, & \text{if } f(\mathbf{x}_i)y_i > 1 \end{cases}$$



Why Hinge Loss?

- Let $\eta(\mathbf{x}) = \mathbb{P}(Y = 1 | \mathbf{X} = \mathbf{x})$. The expected hinge loss: hinge risk.

$$\begin{aligned} R_{\text{hinge}}(f) &= \mathbb{E}_{\mathbf{X}, Y} [L_{\text{hinge}}(f(\mathbf{X}), Y)] \\ &= \mathbb{E}_{\mathbf{X}} \left[\eta(\mathbf{X})(1 - f(\mathbf{X}))_+ + (1 - \eta(\mathbf{X}))(1 + f(\mathbf{X}))_+ \right] \end{aligned}$$

- Suppose that $f(\mathbf{X}) \in [-1, 1]$, for any \mathbf{X} , we have (pls verify in class)

$$\begin{aligned} &\eta(\mathbf{X})(1 - f(\mathbf{X})) + (1 - \eta(\mathbf{X}))(1 + f(\mathbf{X})) \\ &= \eta(\mathbf{X}) - 2\eta(\mathbf{X})f(\mathbf{X}) + 1 + f(\mathbf{X}) - \eta(\mathbf{X}) \\ &= f(\mathbf{X})(1 - 2\eta(\mathbf{X})) + 1. \end{aligned}$$

- The optimal function f_{hinge}^* minimizing $R_{\text{hinge}}(f)$ (why?)
 - If $\eta(\mathbf{X}) < 1/2$, hinge loss is minimized at $f(\mathbf{X}) = -1$
 - If $\eta(\mathbf{X}) > 1/2$, hinge loss is minimized at $f(\mathbf{X}) = 1$

Why Hinge Loss?

- The optimal classifier (i.e., Bayes classifier) of Binary loss is defined as

$$f^*(\mathbf{x}) = \text{sign}(\eta(\mathbf{x}) - 1/2) = \begin{cases} 1 & \text{if } \mathbb{P}(Y = 1 | \mathbf{X} = \mathbf{x}) > 1/2 \\ 0 & \text{if } \mathbb{P}(Y = 1 | \mathbf{X} = \mathbf{x}) < 1/2 \end{cases}$$

- **Observation:**

- (i) f_{hinge}^* is exactly the Bayes classifier defined above;
- (ii) The hinge loss is a convex function, which makes it possible to minimize the training error in practice.


Another surrogate loss - Logistic Loss

↙ 另一个logistic损失

- Definition of Logistic loss:

$$L_{\log}(f(\mathbf{x}_i), y_i) = \log(1 + \exp(-f(\mathbf{x}_i)y_i))$$

Hinge : $\max(0, 1 - f(\mathbf{x}_i) \cdot y_i)$ # $1 - f(\mathbf{x}_i) \cdot y_i > 0$
Logistic: $\log(1 + e^{-f(\mathbf{x}_i) \cdot y_i})$



Why Logistic Loss?

- The logistic risk:

$$\begin{aligned} R_{\log}(f) &= \mathbb{E}_{\mathbf{X}, Y} \left[\log \left(1 + \exp(-f(\mathbf{X})Y) \right) \right] \\ &= \mathbb{E}_{\mathbf{X}} \left[\eta(\mathbf{X}) \log \left(1 + \exp(-f(\mathbf{X})) \right) + (1 - \eta(\mathbf{X})) \log \left(1 + \exp(f(\mathbf{X})) \right) \right] \end{aligned}$$

- Take the derivative with respect to f , **pls verify the following in class**

$$\begin{aligned} & -\eta(\mathbf{X}) \frac{\exp(-f(\mathbf{X}))}{1 + \exp(-f(\mathbf{X}))} + (1 - \eta(\mathbf{X})) \frac{\exp(f(\mathbf{X}))}{1 + \exp(f(\mathbf{X}))} \\ &= -\eta(\mathbf{X}) \frac{1}{1 + \exp(f(\mathbf{X}))} + (1 - \eta(\mathbf{X})) \frac{\exp(f(\mathbf{X}))}{1 + \exp(f(\mathbf{X}))} \\ &= \frac{\exp(f(\mathbf{X}))}{1 + \exp(f(\mathbf{X}))} - \eta(\mathbf{X}) = 0 \iff f_{\log}^*(\mathbf{X}) = \log \frac{\eta(\mathbf{X})}{1 - \eta(\mathbf{X})} \end{aligned}$$

Connection between binary loss and surrogate losses

- The Bayes classifier $f^*(\mathbf{x}) = \text{sign}(\eta(\mathbf{x}) - 1/2)$
- The optimal classifier of Hinge risk $f_{\text{hinge}}^*(\mathbf{x}) = \text{sign}(\eta(\mathbf{x}) - 1/2)$
- The optimal classifier of Logistic risk $f_{\text{log}}^*(\mathbf{x}) = \log \frac{\eta(\mathbf{x})}{1-\eta(\mathbf{x})}$
- **Question:** what is the connection between these optimal classifiers?
- **Answer:** They are consistent in sign in the sense that signs of $f^*, f_{\text{hinge}}^*, f_{\text{log}}^*$ are always the same, e.g., always positive as long as $\eta(\mathbf{x}) > 1/2$.

Logistic Regression – model details

- To estimate $f_{log}^*(\mathbf{x})$, we need to impose an assumption on the form of:

$$\eta(\mathbf{x}) = \mathbb{P}(Y = 1 | \mathbf{x})$$

- In logistic regression, it is often assumed that

$$\eta(\mathbf{x}) = \frac{\exp(\beta_0 + \boldsymbol{\beta}^T \mathbf{x})}{1 + \exp(\beta_0 + \boldsymbol{\beta}^T \mathbf{x})}, \quad (1)$$

where

- $\mathbf{x} = (x_1, \dots, x_p)^T$ is a p -dimensional predictor
- β_0 and $\boldsymbol{\beta} = (\beta_1, \dots, \beta_p)$ are unknown parameters to be estimated
- $\boldsymbol{\beta}^T \mathbf{x} = \sum_{i=1}^p \beta_i x_i$

Rational behind logistic loss: log odds ratio

- By reformulating (1), we have obtained

$$\exp(\beta_0 + \beta^T \mathbf{x}) = \frac{\mathbb{P}(Y = 1 | \mathbf{X} = \mathbf{x})}{1 - \mathbb{P}(Y = 1 | \mathbf{X} = \mathbf{x})} = \frac{\mathbb{P}(Y = 1 | \mathbf{X} = \mathbf{x})}{\mathbb{P}(Y = 0 | \mathbf{X} = \mathbf{x})},$$

where the last term above is the ratio between the conditional probability of $Y = 1$ and that of $Y = 0$ on $\mathbf{X} = \mathbf{x}$, i.e., “odds ratio.”

- In other words, we can claim that the log-odds is assumed to be linear with respect to β :

$$\beta_0 + \beta^T \mathbf{x} = \log \left(\frac{\mathbb{P}(Y = 1 | \mathbf{X} = \mathbf{x})}{\mathbb{P}(Y = 0 | \mathbf{X} = \mathbf{x})} \right)$$

Interpretability: β_i can then be interpreted as the average change in the log-odds ratio given by a one-unit increase in x_i

Maximum likelihood estimation

- Likelihood function $L(\beta_0, \beta)$:

$$L(\beta_0, \beta) = \prod_{i=1}^n \left(\mathbb{P}(Y = 1 | \mathbf{X} = \mathbf{x}) \right)^{y_i} \left(\mathbb{P}(Y = 0 | \mathbf{X} = \mathbf{x}) \right)^{1-y_i}$$

The probability to find x and y given β_0, β
we have to maximize this prob.

- Logarithm of $L(\beta_0, \beta)$ (pls verify in class):

$$\begin{aligned} \log L(\beta_0, \beta) &= \sum_{i=1}^n \left[y_i \log \left(\mathbb{P}(Y = 1 | \mathbf{X} = \mathbf{x}) \right) + (1 - y_i) \log \left(\mathbb{P}(Y = 0 | \mathbf{X} = \mathbf{x}) \right) \right] \\ &= \sum_{i=1}^n \left[y_i (\beta_0 + \beta^T \mathbf{x}) - \log \left(1 + \exp(\beta_0 + \beta^T \mathbf{x}) \right) \right] \end{aligned}$$

Gradient descent/ascent in the computation

Estimate β_0 and β (Gradient Ascent):

$$\beta_0^{(t+1)} \leftarrow \beta_0^{(t)} + \lambda \sum_{i=1}^n \left[y_i - \frac{\exp(\beta_0^{(t)} + \beta^{(t)T} \mathbf{x})}{1 + \exp(\beta_0^{(t)} + \beta^{(t)T} \mathbf{x})} \right]$$
$$\beta^{(t+1)} \leftarrow \beta^{(t)} + \lambda \sum_{i=1}^n \left[y_i - \frac{\exp(\beta_0^{(t)} + \beta^{(t)T} \mathbf{x})}{1 + \exp(\beta_0^{(t)} + \beta^{(t)T} \mathbf{x})} \right] \mathbf{x}_i$$

Now, we are ready to do classification

We have obtained the estimate for β_0 and β , denoted as $\hat{\beta}_0$ and $\hat{\beta}$, based on which we can estimate $P(Y = 1|\mathbf{X})$ as follows:

$$\hat{\eta}(\mathbf{x}) = \frac{\exp(\hat{\beta}_0 + \hat{\beta}^T \mathbf{x})}{1 + \exp(\hat{\beta}_0 + \hat{\beta}^T \mathbf{x})}$$

Then we make predictions by (recall that $\eta(\mathbf{x}) = P(Y = 1|\mathbf{X})$)

$$\hat{f}(\mathbf{x}) = \begin{cases} 1, & \text{if } \hat{\eta}(\mathbf{x}) > 1/2 \\ 0, & \text{if } \hat{\eta}(\mathbf{x}) < 1/2 \end{cases}$$

If $\hat{\eta}(\mathbf{x}) = 1/2$, then just randomly assign a label to it.

Example

Dataset $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^{5000}$, where $\mathbf{x}_i = (x_{i1}, x_{i2}, x_{i3}, x_{i4})$.

- Features are generated from uniform distribution $x_{il} \sim \text{Unif}(0, 2), l = 1, 2, 3, 4$.
- $\beta_0 = 0.5$ and $\boldsymbol{\beta} = (\beta_1, \beta_2, \beta_3, \beta_4)$ with $\beta_i \sim \text{Unif}(-1, 1)$
- Model:

$$Y_i \sim \text{Bernoulli}\left(\frac{\exp(\beta_0 + \boldsymbol{\beta}^T \mathbf{x})}{1 + \exp(\beta_0 + \boldsymbol{\beta}^T \mathbf{x})}\right),$$

which means

$$P(Y_i = 1 | \mathbf{X}) = \frac{\exp(\beta_0 + \boldsymbol{\beta}^T \mathbf{x})}{1 + \exp(\beta_0 + \boldsymbol{\beta}^T \mathbf{x})}$$

Python Codes – data generation

```
import numpy as np
np.random.seed(2)
n,p = 5000,4 # Set training datasize and dimension of features
X = np.random.uniform(-1,1,[n,p]) # Generation of features
beta = np.random.uniform(0,2,4) # Generation of parameters
beta_0 = 0.5 # Set the intercept term to 0.5
logOdd = (X * beta).sum(axis=1)+beta_0 # Log-odds
Prob = np.exp(logOdd)/(1+np.exp(logOdd)) # Probability
Y = np.array(Prob - np.random.uniform(0,1,n)>0,dtype=int) # Generate labels
```




```

Beta_0_hat = 0, # Initialization of intercept term
Beta_hat = np.zeros(p) # Initialization of beta
lamb = 0.1 # Learning rate
Error = [] # Error set
for i in range(2000): # Iterations of gradient ascent
    logOdd_hat = (X * Beta_hat).sum(axis=1) + Beta_0_hat
    Beta_0_hat = Beta_0_hat + lamb * np.mean(Y - np.exp(logOdd_hat) / (1 + np.exp(logOdd_hat)))
    Beta_hat = Beta_hat + lamb * ((Y - np.exp(logOdd_hat) / (1 + np.exp(logOdd_hat))) * X.T).mean(axis=1)
    Error.append(np.linalg.norm(Beta_hat - beta)**2)

import matplotlib.pyplot as plt
plt.plot(np.arange(0, 2000), Error)
plt.xlabel('Number of iterations')
plt.ylabel('$\| \hat{\beta} - \beta \|^2$')
plt.grid()

```

Example: gradient ascent for logistic regression




Over/under-fitting problem

- Overfitting occurs when a model learns the training data too well, capturing noise and making it perform poorly on new, unseen data.
- Underfitting, on the other hand, happens when a model is too simple to capture the underlying patterns in the data, resulting in poor performance on both the training and test data.
- Let's use a simple example with polynomial regression and visualize the above with the out-of-sample (OOS) metrics.


An example - underfitting

Degree 1: Underfitting (Too simple) - The model is not able to capture the underlying pattern in the data.



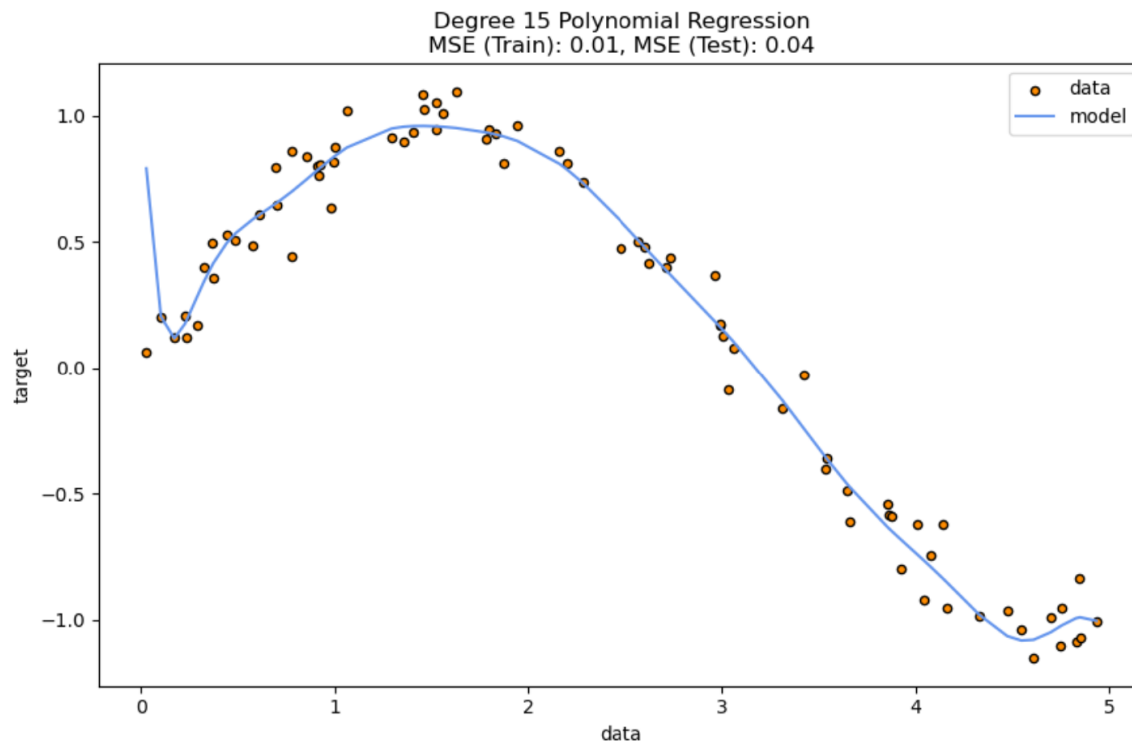
An example - good fit

Degree 5: Good fit - The model captures the underlying pattern well and generalizes to the test data.



An example - overfitting

Degree 15: Overfitting (Too complex) - The model fits the training data too closely, capturing noise and performing poorly on new data.



Confusion matrix – measure the performance of classification

- In Machine Learning, to measure the performance of the classification model, such as logistic regression, we use the confusion matrix.
- A confusion matrix is a matrix that displays the number of accurate and inaccurate classification outcomes for each input instance \mathbf{X} .

Classification outcomes

To measure the performance of classification, we have the following metrics:

- true positives (TP): occurs when the model accurately predicts a positive data point, i.e., $\hat{y} = y = 1$.
- true negatives (TN): occurs when the model accurately predicts a negative data point, i.e., $\hat{y} = y = -1$.
- false positives (FP): occurs when the model predicts a positive data point incorrectly, i.e., $\hat{y} = 1$ but $y = -1$.
- false negatives (FN): occurs when the model predicts a negative data point incorrectly, i.e., $\hat{y} = -1$ but $y = 1$.

An example : dog recognition

Dog: $Y = 1$ & Not Dog: $Y = -1$ (pls verify this table in class!)

index	actual	predicted	Result
1	Dog	Dog	TP
2	Dog	Not Dog	FN
3	Dog	Dog	TP
4	Not Dog	Not Dog	TN
5	Dog	Dog	TP
6	Not Dog	Dog	FP
7	Dog	Dog	TP
8	Dog	Dog	TP
9	Not Dog	Not Dog	TN
10	Not Dog	Not Dog	TN

An example - counts

Pls take a min to count....

- Actual Dog Counts = ?
- Actual Not Dog Counts = ?
- True Positive Counts = ?
- False Positive Counts = ?
- True Negative Counts = ?
- False Negative Counts = ?

An example - counts

- Actual Dog Counts = 6
- Actual Not Dog Counts = 4
- True Positive Counts = 5
- False Positive Counts = 1
- True Negative Counts = 3
- False Negative Counts = 1

An example - construct the confusion matrix

- Construct the Confusion Matrix

		Actual	
		Dog	Not Dog
Predicted	Dog	True Positive (TP =5)	False Positive (FP=1)
	Not Dog	False Negative (FN =1)	True Negative (TN=3)

- How to use the Confusion Matrix for assessing a classification model's performance ?

Classification metrics based on confusion matrix: Accuracy

- Accuracy is used to measure the performance of the model. It is the ratio of total correct instances to the total instances.
- $$\text{Accuracy} = \frac{TP+TN}{TP+TN+FP+FN}$$
- For the above case: Accuracy = ?
- For the above case: $\text{Accuracy} = (5+3)/(5+3+1+1) = 8/10 = 0.8$

Classification metrics based on confusion matrix: Precision

- Precision is a measure of how accurate a model's positive predictions are. Basically, it answers the question “What proportion of positive identifications was actually correct?”
- It is defined as the ratio of true positive predictions to the total number of positive predictions made by the model.
- $\text{Precision} = \frac{TP}{TP+FP}$
- For the above case: Precision = ?
- For the above case: Precision = $5/(5+1) = 5/6 = 0.8333$

Classification metrics based on confusion matrix: Recall

- Recall measures the effectiveness of a classification model in identifying all relevant instances from a dataset. Basically, it answers the question “What proportion of actual positives was identified correctly?”
- It is defined as the ratio of the number of true positive (TP) instances to the sum of true positive and false negative (FN) instances.
- $\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$
- For the above case: $\text{Recall} = ?$
- For the above case: $\text{Recall} = 5/(5+1) = 5/6 = 0.8333$

Classification metrics based on confusion matrix: F-1 Score

- F1-score is used to evaluate the overall performance of a classification model. It is the harmonic mean of precision and recall
- The F1 score is named as such because it is a combination of two other metrics: precision (P) and recall (R). The "F" in F1 stands for "F-measure" or "F-score," and the "1" indicates that it is computed as the harmonic mean of precision and recall.
- The harmonic mean is often used to calculate the average of the ratios or rates.
- The harmonic mean can be expressed as the reciprocal of the arithmetic mean of the reciprocals of the given set of observations.
- For example, harmonic mean of 1, 4, 4 is

$$\left(\frac{1^{-1} + 4^{-1} + 4^{-1}}{3}\right)^{-1} = 2$$

Classification metrics based on confusion matrix: F-1 Score

- **Pls verify this in class:** $F1\text{-score} = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$
- For the above case: F1-Score=?
- For the above case: $F1\text{-Score} = (2 * 0.8333 * 0.8333) / (0.8333 + 0.8333) = 0.8333$

Exercise: Confusion Matrix and Classification Metrics Calculation

Problem statement

- You have a binary classification model used to predict whether an email is spam (positive class) or not spam (negative class). After testing the model on a dataset of 100 emails, you get the following results:
 - 40 emails are correctly identified as spam (True Positives).
 - 10 emails are incorrectly identified as spam (False Positives).
 - 30 emails are correctly identified as not spam (True Negatives).
 - 20 emails are incorrectly identified as not spam (False Negatives).

Exercise: Confusion Matrix and Classification Metrics Calculation

Your task

- Construct a confusion matrix from these results.
- Calculate the following metrics:
 - Accuracy
 - Precision
 - Recall
 - F-1 Score

Exercise: Confusion Matrix and Classification Metrics Calculation

Solution

- confusion matrix:
 - True Positives (TP): 40
 - False Positives (FP): 10
 - True Negatives (TN): 30
 - False Negatives (FN): 20
- Calculate Accuracy: $\text{Accuracy} = \frac{40+30}{40+30+10+20} = 0.7$
- Precision = $\frac{40}{40+10} = 0.8$
- Recall = $\frac{40}{40+20} = \frac{2}{3}$
- F-1 Score = $2 \times \frac{0.8 \times 0.667}{0.8 + 0.667} \approx 0.727$

Application of logistic regression: Banknote authentication dataset

- The Banknote authentication dataset is used for the task of classifying whether a banknote is authentic or not based on certain features extracted from images.
- Variables meaning
 - β_0 : the intercept term;
 - $\beta_1, \beta_2, \beta_3, \beta_4$ are the coefficients associated with each feature.
 - x_1 represents the Variance of the Wavelet Transformed image.
 - x_2 represents the Skewness of the Wavelet Transformed image.
 - x_3 represents the Curtosis of the Wavelet Transformed image.
 - x_4 represents the Entropy of the image.

Application of logistic regression: Banknote authentication dataset

- The logistic regression model makes the following assumption:

$$P(\text{authentic}) = \frac{1}{1 + \exp\{\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_4 X_4\}}$$

- This probability is then used to make a classification decision.
 - If $P(\text{authentic}) \geq 0.5$, the model predicts the banknote as authentic (Class 1)
 - If $P(\text{authentic}) < 0.5$, the model predicts the banknote as not authentic (Class 0)

Application of logistic regression: Results

Pls verify confusion matrix, precision, recall.... by hand after class

```
      Variance  Skewness  Curtosis  Entropy  Class
0    3.62160    8.6661   -2.8073  -0.44699    0
1    4.54590    8.1674   -2.4586  -1.46210    0
2    3.86600   -2.6383    1.9242   0.10645    0
3    3.45660    9.5228   -4.0112  -3.59440    0
4    0.32924   -4.4552    4.5718  -0.98880    0
Accuracy: 0.98
Confusion Matrix:
[[144   4]
 [  2 125]]
Classification Report:
              precision    recall  f1-score   support

         0           0.99       0.97       0.98         148
         1           0.97       0.98       0.98         127

 accuracy              0.98         275
  macro avg           0.98       0.98       0.98         275
  weighted avg           0.98       0.98       0.98         275
```

Statsmodels: Python package

- Statsmodels is primarily focused on statistical modeling and hypothesis testing. It provides tools for estimating and testing various statistical models, including linear regression, logistic regression, time-series analysis, and more.
- Statsmodels includes modules for performing hypothesis tests, constructing confidence intervals, and fitting different types of statistical models with an emphasis on providing detailed statistical information.
- Statsmodels is commonly used in academic research, econometrics, and any scenario where a detailed statistical analysis is required.

Statsmodels: Linear regression example

```
import statsmodels.api as sm
import numpy as np

# Generate some random data for demonstration
np.random.seed(42)
X = np.random.rand(100, 2)
y = 3 * X[:, 0] + 2 * X[:, 1] + 1 + 0.1 * np.random.randn(100)

# Add a constant term to the independent variable
X = sm.add_constant(X)

# Create a linear model
model = sm.OLS(y, X)
results = model.fit()

# Print detailed statistical summary
print(results.summary())
```

Statsmodels: output

```

                        OLS Regression Results
=====
Dep. Variable:          y      R-squared:                0.991
Model:                  OLS    Adj. R-squared:            0.991
Method:                 Least Squares    F-statistic:          5655.
Date:                   Thu, 21 Dec 2023    Prob (F-statistic):    3.86e-101
Time:                   10:50:14    Log-Likelihood:        89.304
No. Observations:       100    AIC:                   -172.6
Df Residuals:           97    BIC:                   -164.8
Df Model:                2
Covariance Type:        nonrobust
=====
                        coef    std err          t      P>|t|      [0.025    0.975]
-----
const                0.9772     0.026    37.651     0.000     0.926     1.029
x1                   3.0339     0.033    91.615     0.000     2.968     3.100
x2                   2.0355     0.035    57.426     0.000     1.965     2.106
=====
Omnibus:                5.986    Durbin-Watson:          2.104
Prob(Omnibus):           0.050    Jarque-Bera (JB):        5.624
Skew:                    0.439    Prob(JB):                0.0601
Kurtosis:                3.761    Cond. No.                 5.22
=====
```

Scikit-learn: Python Lib (a collection of Python packages)

- Purpose: Scikit-learn is a versatile machine learning library that provides tools for various machine learning tasks, including classification, regression, clustering, dimensionality reduction, and more.
- Functionality: Scikit-learn focuses on providing a consistent interface for various machine learning algorithms, making it easy to train models, perform feature engineering, and evaluate model performance.
- Use Cases: Scikit-learn is widely used in industry for building and deploying machine learning models in areas such as image recognition, natural language processing, and predictive analytics.